

[Peer Assessments \(https://class.coursera.org/rprog-005/human\\_grading/\)](https://class.coursera.org/rprog-005/human_grading/)

/ Programming Assignment 2: Lexical Scoping

[Help \(https://class.coursera.org/rprog-005/help/peergrading?url=https%3A%2F%2Fclass.coursera.org%2Frprog-005%2Fhuman\\_grading%2Fview%2Fcourses%2F972576%2Fassessments%2F3%2Fsubmissions\)](https://class.coursera.org/rprog-005/help/peergrading?url=https%3A%2F%2Fclass.coursera.org%2Frprog-005%2Fhuman_grading%2Fview%2Fcourses%2F972576%2Fassessments%2F3%2Fsubmissions)

due in 1day 19h

#### Submission Phase

1. Do assignment ☐ (/rprog-005/human\_grading/view/courses/972576/assessments/3/submissions)

#### Evaluation Phase

2. Evaluate peers  (/rprog-005/human\_grading/view/courses/972576/assessments/3/peerGradingSets)

#### Results Phase

3. See results  (/rprog-005/human\_grading/view/courses/972576/assessments/3/results/mine)

☐ In accordance with the Honor Code, I certify that my answers here are my own work, and that I have appropriately acknowledged all external sources (if any) that were used in this work.

[Save draft](#)

[Submit for grading](#)

## Introduction

This second programming assignment will require you to write an R function is able to cache potentially time-consuming computations. For example, taking the mean of a numeric vector is typically a fast operation. However, for a very long vector, it may take too long to compute the mean, especially if it has to be computed repeatedly (e.g. in a loop). If the contents of a vector are not changing, it may make sense to cache the value of the mean so that when we need it again, it can be looked up in the cache rather than recomputed. In this Programming Assignment will take advantage of the scoping rules of the R language and how they can be manipulated to preserve state inside of an R object.

### Example: Caching the Mean of a Vector

In this example we introduce the `<<-` operator which can be used to assign a value to an object in an environment that is different from the current environment. Below are two functions that are used to create a special object that stores a numeric vector and cache's its mean.

The first function, `makeVector` creates a special "vector", which is really a list containing a function to

1. set the value of the vector
2. get the value of the vector
3. set the value of the mean
4. get the value of the mean

```
makeVector <- function(x = numeric()) {  
  m <- NULL  
  set <- function(y) {  
    x <- y  
    m <- NULL  
  }  
  get <- function() x  
  setmean <- function(mean) m <- mean  
  getmean <- function() m  
  list(set = set, get = get,  
        setmean = setmean,  
        getmean = getmean)  
}
```

The following function calculates the mean of the special "vector" created with the above function.

However, it first checks to see if the mean has already been calculated. If so, it gets the mean from the cache and skips the computation. Otherwise, it calculates the mean of the data and sets the value of the mean in the cache via the `setmean` function.

```
cachemean <- function(x, ...) {  
  m <- x$getmean()  
  if(!is.null(m)) {  
    message("getting cached data")  
    return(m)  
  }  
  data <- x$get()  
  m <- mean(data, ...)  
  x$setmean(m)  
  m  
}
```

## Assignment: Caching the Inverse of a Matrix

Matrix inversion is usually a costly computation and there may be some benefit to caching the inverse of a matrix rather than compute it repeatedly (there are also alternatives to matrix inversion that we will not discuss here). Your assignment is to write a pair of functions that cache the inverse of a matrix.

Write the following functions:

1. `makeCacheMatrix`: This function creates a special "matrix" object that can cache its inverse.
2. `cacheSolve`: This function computes the inverse of the special "matrix" returned by `makeCacheMatrix` above. If the inverse has already been calculated (and the matrix has not changed), then the `cacheSolve` should retrieve the inverse from the cache.

Computing the inverse of a square matrix can be done with the solve function in R. For example, if X is a square invertible matrix, then solve(X) returns its inverse.

For this assignment, assume that the matrix supplied is always invertible.

In order to complete this assignment, you must do the following:

1. Fork the GitHub repository containing the stub R files at <https://github.com/rdpeng/ProgrammingAssignment2> (<https://github.com/rdpeng/ProgrammingAssignment2>) to create a copy under your own account.
2. Clone your forked GitHub repository to your computer so that you can edit the files locally on your own machine.
3. Edit the R file contained in the git repository and place your solution in that file (please do not rename the file).
4. Commit your completed R file into YOUR git repository and push your git branch to the GitHub repository under your account.
5. Submit to Coursera the URL to your GitHub repository that contains the completed R code for the assignment.

In addition to submitting the URL for your GitHub repository, you will need to submit the **40 character SHA-1 hash** (as string of numbers from 0-9 and letters from a-f) that identifies the repository commit that contains the version of the files you want to submit. You can do this in GitHub by doing the following

1. Going to your GitHub repository web page for this assignment
2. Click on the “?? commits” link where ?? is the number of commits you have in the repository. For example, if you made a total of 10 commits to this repository, the link should say “10 commits”.
3. You will see a list of commits that you have made to this repository. The most recent commit is at the very top. If this represents the version of the files you want to submit, then just click the “copy to clipboard” button on the right hand side that should appear when you hover over the SHA-1 hash. Paste this SHA-1 hash into the course web site when you submit your assignment. If you don't want to use the most recent commit, then go down and find the commit you want and copy the SHA-1 hash.

A valid submission will look something like (this is just an **example!**)

<https://github.com/rdpeng/ProgrammingAssignment2> (<https://github.com/rdpeng/ProgrammingAssignment2>)

7c376cc5447f11537f8740af8e07d6facc3d9645

## Grading

This assignment will be graded via peer assessment.

Please submit the URL pointing to your GitHub repository containing the completed R code for this assignment. (NOTE: Please submit the URL for the entire repository, not for any individual files.)

On a **separate line**, submit the SHA-1 hash identifier corresponding to the commit that contains the full submission (the identifier should be a string of 40 characters containing numbers 0-9 and letters a-f).

<b>B</b>	<i>I</i>			 Link	<code>&lt;code&gt;</code>	Math		Edit: Rich ▼	Preview
----------	----------	---	---	--	---------------------------	------	--	--------------	---------

## Evaluation/feedback on the above work

**Note:** this section can only be filled out during the evaluation phase.

Was a valid GitHub URL containing a git repository submitted?

0 points: A valid GitHub URL was NOT submitted (or URL is broken)

1 point: The submitted URL points to a GitHub repository

Does the GitHub repository contain at least one commit beyond the original fork?

0 points: No, there are no commits beyond the original fork

1 point: Yes, there is at least one commit beyond the original fork

Was a SHA-1 submitted indicating a specific commit in the GitHub repository?

0 points: No, there was no SHA-1 submitted

1 point: Yes, a SHA-1 was submitted

If a SHA-1 was submitted along with the GitHub URL, does the SHA-1 correspond to a specific commit in the repository?

0 points: No, there is no commit in the repository corresponding to that SHA-1  
2 points: Yes, there is a commit in the repository corresponding to that SHA-1

## Overall evaluation/feedback

**Note:** this section can only be filled out during the evaluation phase.

Does the GitHub repository contain an R file containing code implementing the completed assignment?

NOTE: There is no need to run the code here, but rather you should visually inspect the R file in the GitHub repository and check to see that there is code there beyond the original stub that was committed there by the instructor. Do not attempt to judge whether the code is correct or not.

0 points: No, the R file does not contain any code implementing the completed assignment  
1 point: The R file contains code implementing a partially completed assignment  
2 points: The R file contains code implementing the complete assignment

Does the R file containing the code have any comments explaining what the code does?

NOTE: The makeCacheMatrix and cacheSolve functions should both be documented with explanatory comments. There maybe other functions in the R file but they do not need to be commented.

0 points: There are no explanatory comments in the R file  
1 points: One of the functions has some corresponding explanatory comments in the R file  
2 points: Both functions have corresponding explanatory comments in the R file

Does the R code implementing the 'makeCacheMatrix' function appear to be correct, to the best of your ability to judge?

NOTE: Do not run the R code on your own computer. Please examine the R code and determine to the best of your ability, whether the solution presented appears to match the requirements of the assignment. A correct solution (as best you can determine) gets 1 point and a solution that is well-written and easy to read gets an additional 1 point.

0 points: The solution contains at least one identifiable problem

1 point: The solution appears to be correct

2 points: The solution appears to be correct and is written in a well-formatted and easily readable style

Does the R code implementing the 'cacheSolve' function appear to be correct, to the best of your ability to judge?

NOTE: Do not run the R code on your own computer. Please examine the R code and determine to the best of your ability, whether the solution presented appears to match the requirements of the assignment. A correct solution (as best you can determine) gets 1 point and a solution that is well-written and easy to read gets an additional 1 point.

0 points: The solution contains at least one identifiable problem

1 point: The solution appears to be correct

2 points: The solution appears to be correct and is written in a well-formatted and easily readable style

Use this space to provide constructive feedback to the student who submitted the work. Point out both strengths and weaknesses in the submission and provide advice about how the work could be improved in the future.

You've written 0 words

---

☐ In accordance with the Honor Code, I certify that my answers here are my own work, and that I have appropriately acknowledged all external sources (if any) that were used in this work.

[Save draft](#)

[Submit for grading](#)